

## Current Industry Practices and the Role of Logistic Regression in Mental Health

Logistic regression is a widely used machine learning technique in the healthcare industry for predicting binary outcomes, such as the presence or absence of depression. It is favored for its simplicity, interpretability, and effectiveness in handling structured clinical data. In practice, logistic regression models utilize patient demographics, lifestyle factors, medical history, and psychological assessments to predict depression risk, supporting early diagnosis and personalized treatment plans in clinical settings.

## Research Landscape and Recent Developments in Depression Prediction

Recent studies demonstrate the effectiveness of logistic regression models in accurately predicting depression by identifying significant risk factors such as age, sleep patterns, and physical activity. For instance, research by Smith et al. (2022) and Zhang et al. (2023) showed logistic regression models achieving high accuracy in depression classification. These models are praised for their ability to provide clear, interpretable results that can be easily communicated to healthcare professionals for clinical decision-making.

## Challenges, Limitations, and Future Directions

While logistic regression is a robust tool for depression prediction, challenges such as handling multicollinearity, feature scaling, and dealing with missing data can affect model performance. Recent literature explores techniques like regularization and data imputation to overcome these limitations. The field is also moving towards integrating logistic regression with advanced machine learning methods to enhance predictive accuracy while maintaining interpretability, addressing the evolving needs of precision mental health care.

|    | Variable Name                | Type        | Description  | Values   |
|----|------------------------------|-------------|--|--|
| 0  | Age                          | Numerical   | The age of the individual.                                   | Numbers  |
| 1  | Marital Status               | Categorical | The marital status of the individual.                        | Married, Single, Widowed, Divorced                                     |
| 2  | Education Level              | Categorical | The highest level of education attained by the individual.   | Bachelor's Degree, High School, Associate Degree, Master's Degree, PhD |
| 3  | Number of Children           | Numerical   | The number of children the individual has.                   | 0, 1, 2, 3, 4  |
| 4  | Smoking Status               | Categorical | Whether the individual smokes.                               | Non-smoker, Former, Current  |
| 5  | Physical Activity Level      | Categorical | The level of physical activity engaged in by the individual. | Sedentary, Moderate, Active  |
| 6  | Employment Status            | Categorical | The current employment status of the individual.             | Employed, Unemployed   |
| 7  | Income                       | Numerical   | The annual income of the individual.                         | Numbers  |
| 8  | Alcohol Consumption          | Categorical | The level of alcohol consumption by the individual.          | Moderate, Low, High  |
| 9  | Dietary Habits               | Categorical | The eating habits of the individual.                         | Unhealthy, Moderate, Healthy   |
| 10 | Sleep Patterns               | Categorical | The sleep patterns of the individual.                        | Fair, Poor, Good   |
| 11 | History of Mental Illness    | Categorical | Whether the individual has a history of mental illness.      | No, Yes  |
| 12 | History of Substance Abuse   | Categorical | Whether the individual has a history of substance abuse.     | No, Yes  |
| 13 | Family History of Depression | Categorical | Whether there is a family history of depression.             | No, Yes  |
| 14 | Depression                   | Categorical | Whether the individual has any chronic medical conditions.   | No, Yes  |
| 15 | Name                         | Object      | Name of Individual   | —  |

## Libraries

```
In [55]: import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import OneHotEncoder

import time

In [2]: import pyarrow as pa
import pyarrow.parquet as pq

file = pd.read_csv("depression_data.csv")
table = pa.Table.from_pandas(file)
pq.write_table(table, "depression_data.parquet")
```

## Tell Time - Decorator for calculating the time

```
In [3]: def tell_time(function, *args, **kwargs):
def wrapper(*args, **kwargs):
start = time.time()
done = function(*args, **kwargs)
print(f"function.__name__() function took - {(time.time()-start)//60} Mins {(time.time()-start)%60:.3f} Sec")
return done
return wrapper

In [4]: @tell_time
def read_csv_data(file):
return pd.read_csv(file)

In [5]: @tell_time
def read_parquet_data(file):
return pd.read_parquet(file)
```

## Reading the Dataset

```
In [6]: csv_df = read_csv_data("depression_data.csv")
read_csv_data() function took - 0.0 Mins 0.785 Sec

In [7]: parquet_df = read_parquet_data("depression_data.parquet")
read_parquet_data() function took - 0.0 Mins 0.463 Sec

In [8]: df = parquet_df

In [9]: df.rename(columns={"Chronic Medical Conditions": "Depression"}, inplace=True)

In [10]: df.head(3)
```

|   | Name             | Age | Marital Status | Education Level   | Number of Children | Smoking Status | Physical Activity Level | Employment Status | Income    | Alcohol Consumption | Dietary Habits | Sleep Patterns | History of Mental Illness | History of Substance Abuse | Family History of Depression | Depression |
|---|------------------|-----|----------------|-------------------|--------------------|----------------|-------------------------|-------------------|-----------|---------------------|----------------|----------------|---------------------------|----------------------------|------------------------------|------------|
| 0 | Christine Barker | 31  | Married        | Bachelor's Degree | 2                  | Non-smoker     | Active                  | Unemployed        | 26265.67  | Moderate            | Moderate       | Fair           | Yes                       | No                         | Yes                          | Yes        |
| 1 | Jacqueline Lewis | 55  | Married        | High School       | 1                  | Non-smoker     | Sedentary               | Employed          | 42710.36  | High                | Unhealthy      | Fair           | Yes                       | No                         | No                           | Yes        |
| 2 | Shannon Church   | 78  | Widowed        | Master's Degree   | 1                  | Non-smoker     | Sedentary               | Employed          | 125332.79 | Low                 | Unhealthy      | Good           | No                        | No                         | Yes                          | No         |

## Inspection of Dataset

```
In [11]: df.shape

Out[11]: (413768, 16)

In [12]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 413768 entries, 0 to 413767
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                ---
0   Name                                413768 non-null object
1   Age                                413768 non-null int64
2   Marital Status                     413768 non-null object
3   Education Level                    413768 non-null object
4   Number of Children                 413768 non-null int64
5   Smoking Status                     413768 non-null object
6   Physical Activity Level             413768 non-null object
7   Employment Status                  413768 non-null object
8   Income                             413768 non-null float64
9   Alcohol Consumption                413768 non-null object
10  Dietary Habits                     413768 non-null object
11  Sleep Patterns                     413768 non-null object
12  History of Mental Illness           413768 non-null object
13  History of Substance Abuse          413768 non-null object
14  Family History of Depression        413768 non-null object
15  Depression                          413768 non-null object
dtypes: float64(1), int64(2), object(13)
memory usage: 50.5+ MB

In [13]: df["Number of Children"].value_counts()

Out[13]:
Number of Children
0    155232
2     83961
1     83925
3     76974
4     13676
Name: count, dtype: int64

In [14]: df["Number of Children"] = df["Number of Children"].astype("object")
```

## Statistical Inferences - 5 Point Summary & describe

```
In [15]: df.describe(include="number")

Out[15]:
```

|       | Age           | Income        |
|-------|---------------|---------------|
| count | 413768.000000 | 413768.000000 |
| mean  | 49.000713     | 50661.707971  |
| std   | 18.158759     | 40624.100565  |
| min   | 18.000000     | 0.410000      |
| 25%   | 33.000000     | 21001.030000  |
| 50%   | 49.000000     | 37520.135000  |
| 75%   | 65.000000     | 76616.300000  |
| max   | 80.000000     | 209995.220000 |

```
In [16]: df.describe(exclude="number")

Out[16]:
```

|        | Name          | Marital Status | Education Level   | Number of Children | Smoking Status | Physical Activity Level | Employment Status | Alcohol Consumption | Dietary Habits | Sleep Patterns | History of Mental Illness | History of Substance Abuse | Family History of Depression | Depression |
|--------|---------------|----------------|-------------------|--------------------|----------------|-------------------------|-------------------|---------------------|----------------|----------------|---------------------------|----------------------------|------------------------------|------------|
| count  | 413768        | 413768         | 413768            | 413768             | 413768         | 413768                  | 413768            | 413768              | 413768         | 413768         | 413768                    | 413768                     | 413768                       | 413768     |
| unique | 196851        | 4              | 5                 | 5                  | 3              | 3                       | 2                 | 3                   | 3              | 3              | 2                         | 2                          | 2                            | 2          |
| top    | Michael Smith | Married        | Bachelor's Degree | 0                  | Non-smoker     | Sedentary               | Employed          | Moderate            | Unhealthy      | Fair           | No                        | No                         | No                           | No         |
| freq   | 198           | 240444         | 124329            | 155232             | 247416         | 176850                  | 265659            | 173440              | 170817         | 196789         | 287943                    | 284880                     | 302515                       | 277561     |

## Cleaning the Data - Nulls & Duplicates

```
In [17]: df.duplicated().sum()
```

```
Out[17]: np.int64(0)
```

```
In [18]: df.isnull().sum()
```

```
Out[18]: Name                0
Age                0
Marital Status     0
Education Level    0
Number of Children 0
Smoking Status     0
Physical Activity Level 0
Employment Status  0
Income            0
Alcohol Consumption 0
Dietary Habits     0
Sleep Patterns     0
History of Mental Illness 0
History of Substance Abuse 0
Family History of Depression 0
Depression         0
dtype: int64
```

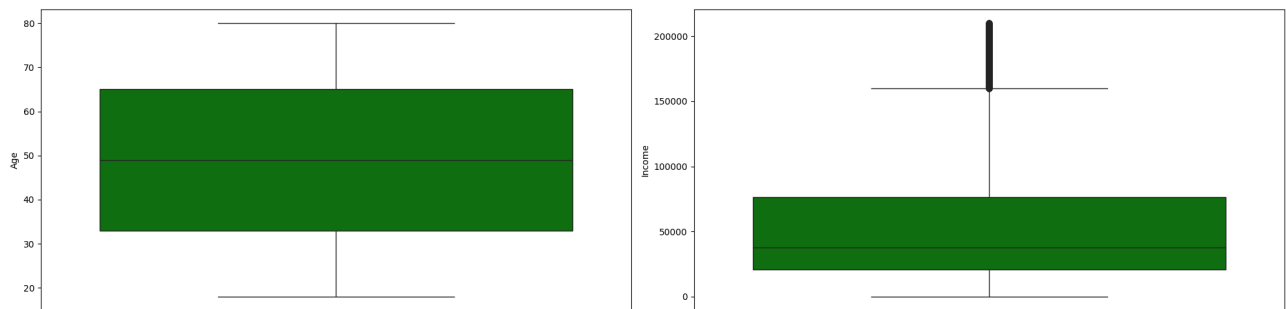
```
In [19]: df.drop("Name", axis=1, inplace=True)
```

## Plots

```
In [20]: fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20,5))

for feature, ax_object in zip(df.select_dtypes(include="number").columns, ax.flatten()):
    sns.boxplot(df[feature], ax=ax_object, color="green")

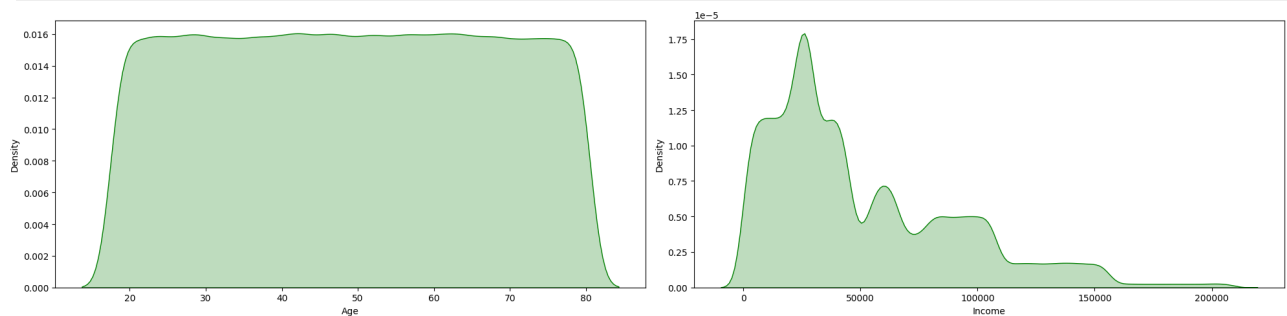
plt.tight_layout()
plt.show()
```



```
In [25]: fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20,5))

for feature, ax_object in zip(x_train.select_dtypes(include="number").columns, ax.flatten()):
    sns.kdeplot(x_train[feature], ax=ax_object, color="green", fill=True)

plt.tight_layout()
plt.show()
```



```
In [21]: fig, axes = plt.subplots(nrows=4, ncols=4, figsize=(20, 7))

for i in range(13):
    feature = ['Marital Status', 'Education Level', 'Number of Children',
               'Smoking Status', 'Physical Activity Level', 'Employment Status',
               'Alcohol Consumption', 'Dietary Habits', 'Sleep Patterns',
               'History of Mental Illness', 'History of Substance Abuse',
               'Family History of Depression', 'Depression']

    category = df[feature[i]].value_counts().index
    percentage = [f"{i:.2f}%" for i in (df[feature[i]].value_counts().values/df.shape[0])*100]

    f_table = pd.DataFrame({"Category":category, "Percentage":percentage})

    axes[i//4, i%4].axis('tight')
    axes[i//4, i%4].axis('off')
    axes[i//4, i%4].table(cellText=f_table.values, colLabels=f_table.columns, loc="center", cellLoc="center", fontsize=15, colColours=["yellow"]*13)
    axes[i//4, i%4].set_title(feature[i])

axes[3, 1].axis('off')
axes[3, 2].axis('off')
axes[3, 3].axis('off')
```

```
plt.tight_layout()
plt.show()
```

| Marital Status |            | Education Level   |            | Number of Children |            | Smoking Status |            |
|----------------|------------|-------------------|------------|--------------------|------------|----------------|------------|
| Category       | Percentage | Category          | Percentage | Category           | Percentage | Category       | Percentage |
| Married        | 58.11%     | Bachelor's Degree | 30.05%     | 0                  | 37.52%     | Non-smoker     | 59.80%     |
| Single         | 17.43%     | High School       | 28.74%     | 2                  | 20.29%     | Former         | 28.08%     |
| Widowed        | 16.55%     | Associate Degree  | 19.33%     | 1                  | 20.28%     | Current        | 12.12%     |
| Divorced       | 7.91%      | Master's Degree   | 17.83%     | 3                  | 18.60%     |                |            |
|                |            | PhD               | 4.05%      | 4                  | 3.31%      |                |            |

| Physical Activity Level |            | Employment Status |            | Alcohol Consumption |            | Dietary Habits |            |
|-------------------------|------------|-------------------|------------|---------------------|------------|----------------|------------|
| Category                | Percentage | Category          | Percentage | Category            | Percentage | Category       | Percentage |
| Sedentary               | 42.74%     | Employed          | 64.20%     | Moderate            | 41.92%     | Unhealthy      | 41.28%     |
| Moderate                | 38.19%     | Unemployed        | 35.80%     | Low                 | 33.65%     | Moderate       | 41.19%     |
| Active                  | 19.07%     |                   |            | High                | 24.43%     | Healthy        | 17.52%     |

| Sleep Patterns |            | History of Mental Illness |            | History of Substance Abuse |            | Family History of Depression |            |
|----------------|------------|---------------------------|------------|----------------------------|------------|------------------------------|------------|
| Category       | Percentage | Category                  | Percentage | Category                   | Percentage | Category                     | Percentage |
| Fair           | 47.56%     | No                        | 69.59%     | No                         | 68.85%     | No                           | 73.11%     |
| Poor           | 31.32%     | Yes                       | 30.41%     | Yes                        | 31.15%     | Yes                          | 26.89%     |
| Good           | 21.12%     |                           |            |                            |            |                              |            |

| Depression |            |
|------------|------------|
| Category   | Percentage |
| No         | 67.08%     |
| Yes        | 32.92%     |

## Splitting the data

```
In [22]: le_object = LabelEncoder()
df["Depression"] = le_object.fit_transform(df["Depression"])
target = df["Depression"]
df = df.drop(columns="Depression")

In [23]: x_train, x_test, y_train, y_test = train_test_split(df, target, test_size=0.2, random_state=5)

In [24]: print(f'''
x_train shape : {x_train.shape}
x_test shape : {x_test.shape}
y_train shape : {y_train.shape}
y_test shape : {y_test.shape}
''')

x_train shape : (331014, 14)
x_test shape : (82754, 14)
y_train shape : (331014,)
y_test shape : (82754,)
```

## Transformation

### Scaling

```
In [26]: scaler_objects = {}

for i in x_train.select_dtypes(include="number").columns:
    scaler_objects[i] = MinMaxScaler()
    scaler_objects[i].fit(x_train[[i]])
    x_train[i] = scaler_objects[i].transform(x_train[[i]])

In [28]: for i in x_test.select_dtypes(include="number").columns:
x_test[i] = scaler_objects[i].transform(x_test[[i]])
```

### Encoding

```
In [31]: new_df = x_train.select_dtypes(include="number")
new_df.reset_index(inplace=True)
encoder_objects = {}

for i in x_train.select_dtypes(exclude="number").columns:
    # print(i)
    encoder_objects[i] = OneHotEncoder(dtype='int', drop="first")
    dummy_df = encoder_objects[i].fit_transform(x_train[[i]]).toarray()
    new_df = pd.concat([new_df, pd.DataFrame(dummy_df, columns=encoder_objects[i].get_feature_names_out()), axis=1])
    # print(new_df)

x_train = new_df

In [32]: new_df = x_test.select_dtypes(include="number")
new_df.reset_index(inplace=True)

for i in x_test.select_dtypes(exclude="number").columns:
    dummy_df = encoder_objects[i].transform(x_test[[i]]).toarray()
    new_df = pd.concat([new_df, pd.DataFrame(dummy_df, columns=encoder_objects[i].get_feature_names_out()), axis=1])
    # print(new_df)

x_test = new_df

In [37]: x_train.drop(columns="index", inplace=True)
x_test.drop(columns="index", inplace=True)
```

## Base Model - Logistic Regression

```
In [52]: model = LogisticRegression()  
model.fit(x_train, y_train)
```

```
Out[52]: LogisticRegression  
LogisticRegression()
```

```
In [53]: y_pred = model.predict(x_test)
```

```
In [60]: np.unique(y_pred, return_counts=True)
```

```
Out[60]: (array([0]), array([82754]))
```

```
In [56]: print(classification_report(y_test, y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.67      | 1.00   | 0.80     | 55526   |
| 1            | 0.00      | 0.00   | 0.00     | 27228   |
| accuracy     |           |        | 0.67     | 82754   |
| macro avg    | 0.34      | 0.50   | 0.40     | 82754   |
| weighted avg | 0.45      | 0.67   | 0.54     | 82754   |

```
In [ ]:
```