

## Current Industry Practices and the Role of Logistic Regression in Mental Health

Logistic regression is a widely used machine learning technique in the healthcare industry for predicting binary outcomes, such as the presence or absence of depression. It is favored for its simplicity, interpretability, and effectiveness in handling structured clinical data. In practice, logistic regression models utilize patient demographics, lifestyle factors, medical history, and psychological assessments to predict depression risk, supporting early diagnosis and personalized treatment plans in clinical settings.

## Research Landscape and Recent Developments in Depression Prediction

Recent studies demonstrate the effectiveness of logistic regression models in accurately predicting depression by identifying significant risk factors such as age, sleep patterns, and physical activity. For instance, research by Smith et al. (2022) and Zhang et al. (2023) showed logistic regression models achieving high accuracy in depression classification. These models are praised for their ability to provide clear, interpretable results that can be easily communicated to healthcare professionals for clinical decision-making.

## Challenges, Limitations, and Future Directions

While logistic regression is a robust tool for depression prediction, challenges such as handling multicollinearity, feature scaling, and dealing with missing data can affect model performance. Recent literature explores techniques like regularization and data imputation to overcome these limitations. The field is also moving towards integrating logistic regression with advanced machine learning methods to enhance predictive accuracy while maintaining interpretability, addressing the evolving needs of precision mental health care.

	Variable Name	Type	Description	Values
0	Age	Numerical	The age of the individual.	Numbers
1	Marital Status	Categorical	The marital status of the individual.	Married, Single, Widowed, Divorced
2	Education Level	Categorical	The highest level of education attained by the individual.	Bachelor's Degree, High School, Associate Degree, Master's Degree, PhD
3	Number of Children	Numerical	The number of children the individual has.	0, 1, 2, 3, 4
4	Smoking Status	Categorical	Whether the individual smokes.	Non-smoker, Former, Current
5	Physical Activity Level	Categorical	The level of physical activity engaged in by the individual.	Sedentary, Moderate, Active
6	Employment Status	Categorical	The current employment status of the individual.	Employed, Unemployed
7	Income	Numerical	The annual income of the individual.	Numbers
8	Alcohol Consumption	Categorical	The level of alcohol consumption by the individual.	Moderate, Low, High
9	Dietary Habits	Categorical	The eating habits of the individual.	Unhealthy, Moderate, Healthy
10	Sleep Patterns	Categorical	The sleep patterns of the individual.	Fair, Poor, Good
11	History of Mental Illness	Categorical	Whether the individual has a history of mental illness.	No, Yes
12	History of Substance Abuse	Categorical	Whether the individual has a history of substance abuse.	No, Yes
13	Family History of Depression	Categorical	Whether there is a family history of depression.	No, Yes
14	Depression	Categorical	Whether the individual has any chronic medical conditions.	No, Yes
15	Name	Object	Name of Individual	—

## Libraries

```
In [39]: import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression, RidgeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, precision_recall_fscore_support, accuracy_score, confusion_matrix, roc_curve, roc_auc_score
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import VotingClassifier
from matplotlib.colors import ListedColormap

import time
```

```
In [40]: import pyarrow as pa
import pyarrow.parquet as pq

file = pd.read_csv("depression_data.csv")
table = pa.Table.from_pandas(file)
pq.write_table(table, "depression_data.parquet")
```

## Tell Time - Decorator for calculating the time

```
In [41]: def tell_time(function, *args, **kwargs):
def wrapper(*args, **kwargs):
start = time.time()
done = function(*args, **kwargs)
print(f"function {function.__name__}() function took - {(time.time()-start)//60} Mins {(time.time()-start)%60:.3f} Sec")
return done
return wrapper
```

```
In [42]: @tell_time
def read_csv_data(file):
return pd.read_csv(file)
```

```
In [43]: @tell_time
def read_parquet_data(file):
return pd.read_parquet(file)
```

Reading the Dataset

```
In [44]: csv_df = read_csv_data("depression_data.csv")
read_csv_data() function took - 0.0 Mins 0.426 Sec

In [45]: parquet_df = read_parquet_data("depression_data.parquet")
read_parquet_data() function took - 0.0 Mins 0.226 Sec

In [46]: df = parquet_df

In [47]: df.rename(columns={"Chronic Medical Conditions": "Depression"}, inplace=True)

In [48]: df.head(3)
```

Out[48]:

	Name	Age	Marital Status	Education Level	Number of Children	Smoking Status	Physical Activity Level	Employment Status	Income	Alcohol Consumption	Dietary Habits	Sleep Patterns	History of Mental Illness	History of Substance Abuse	Family History of Depression	Depression
0	Christine Barker	31	Married	Bachelor's Degree	2	Non-smoker	Active	Unemployed	26265.67	Moderate	Moderate	Fair	Yes	No	Yes	Yes
1	Jacqueline Lewis	55	Married	High School	1	Non-smoker	Sedentary	Employed	42710.36	High	Unhealthy	Fair	Yes	No	No	Yes
2	Shannon Church	78	Widowed	Master's Degree	1	Non-smoker	Sedentary	Employed	125332.79	Low	Unhealthy	Good	No	No	Yes	No

Inspection of Dataset

```
In [49]: df.shape

Out[49]: (413768, 16)

In [50]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 413768 entries, 0 to 413767
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Name                                413768 non-null object
1   Age                                413768 non-null int64
2   Marital Status                      413768 non-null object
3   Education Level                     413768 non-null object
4   Number of Children                  413768 non-null int64
5   Smoking Status                      413768 non-null object
6   Physical Activity Level              413768 non-null object
7   Employment Status                   413768 non-null object
8   Income                              413768 non-null float64
9   Alcohol Consumption                 413768 non-null object
10  Dietary Habits                      413768 non-null object
11  Sleep Patterns                      413768 non-null object
12  History of Mental Illness            413768 non-null object
13  History of Substance Abuse           413768 non-null object
14  Family History of Depression         413768 non-null object
15  Depression                           413768 non-null object
dtypes: float64(1), int64(2), object(13)
memory usage: 50.5+ MB

In [51]: df["Number of Children"].value_counts()

Out[51]:
Number of Children
0      155232
2       83961
1       83925
3       76974
4       13676
Name: count, dtype: int64

In [52]: df["Number of Children"] = df["Number of Children"].astype("object")
```

Statistical Inferences - 5 Point Summary & describe

```
In [53]: df.describe(include="number")

Out[53]:
```

	Age	Income
count	413768.000000	413768.000000
mean	49.000713	50661.707971
std	18.158759	40624.100565
min	18.000000	0.410000
25%	33.000000	21001.030000
50%	49.000000	37520.135000
75%	65.000000	76616.300000
max	80.000000	209995.220000

```
In [54]: df.describe(exclude="number")
```

Out[54]:

	Name	Marital Status	Education Level	Number of Children	Smoking Status	Physical Activity Level	Employment Status	Alcohol Consumption	Dietary Habits	Sleep Patterns	History of Mental Illness	History of Substance Abuse	Family History of Depression	Depression
count	413768	413768	413768	413768	413768	413768	413768	413768	413768	413768	413768	413768	413768	413768
unique	196851	4	5	5	3	3	2	3	3	3	2	2	2	2
top	Michael Smith	Married	Bachelor's Degree	0	Non-smoker	Sedentary	Employed	Moderate	Unhealthy	Fair	No	No	No	No
freq	198	240444	124329	155232	247416	176850	265659	173440	170817	196789	287943	284880	302515	277561

## Cleaning the Data - Nulls & Duplicates

In [55]: `df.duplicated().sum()`

Out[55]: `np.int64(0)`

In [56]: `df.isnull().sum()`

Out[56]:

Name	0
Age	0
Marital Status	0
Education Level	0
Number of Children	0
Smoking Status	0
Physical Activity Level	0
Employment Status	0
Income	0
Alcohol Consumption	0
Dietary Habits	0
Sleep Patterns	0
History of Mental Illness	0
History of Substance Abuse	0
Family History of Depression	0
Depression	0


dtype: int64

In [57]: `df.drop("Name", axis=1, inplace=True)`

## Plots


In [58]:

```
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20,5))  
  
for feature, ax_object in zip(df.select_dtypes(include="number").columns, ax.flatten()):  
    sns.boxplot(df[feature], ax=ax_object, color="#05C7F2")  
  
plt.tight_layout()  
plt.show()
```



In [59]:

```
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20,5))  
  
for feature, ax_object in zip(df.select_dtypes(include="number").columns, ax.flatten()):  
    sns.kdeplot(df[feature], ax=ax_object, fill=True, color="#05C7F2")  
  
plt.tight_layout()  
plt.show()
```



In [60]:

```
fig, axes = plt.subplots(nrows=4, ncols=4, figsize=(20, 7))  
  
for i in range(13):  
    feature = ['Marital Status', 'Education Level', 'Number of Children',  
              'Smoking Status', 'Physical Activity Level', 'Employment Status',  
              'Alcohol Consumption', 'Dietary Habits', 'Sleep Patterns',  
              'History of Mental Illness', 'History of Substance Abuse',  
              'Family History of Depression', 'Depression']
```

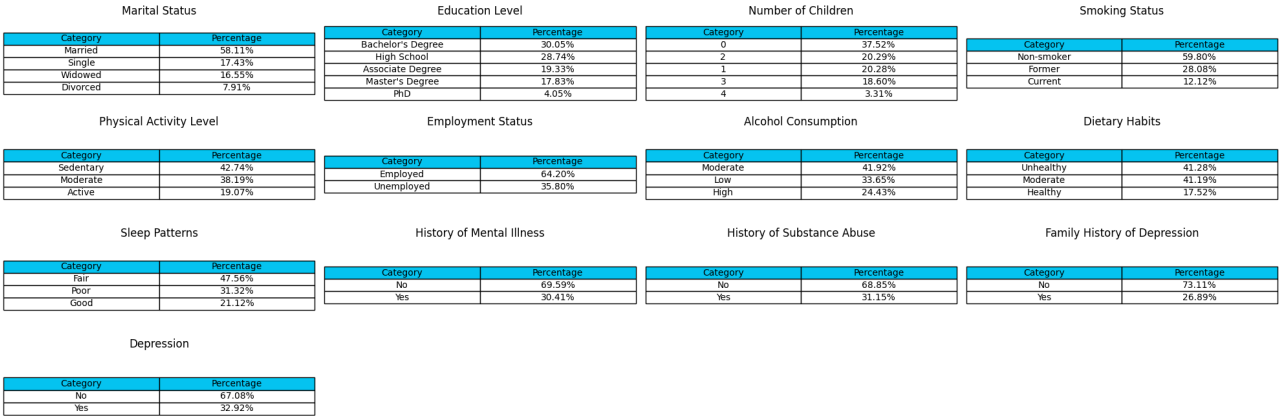
```
category = df[feature[i]].value_counts().index
percentage = [f"{i:.2f}%" for i in (df[feature[i]].value_counts().values/df.shape[0])*100]

f_table = pd.DataFrame({"Category":category, "Percentage":percentage})

axes[i//4, i%4].axis('tight')
axes[i//4, i%4].axis('off')
axes[i//4, i%4].table(cellText=f_table.values, colLabels=f_table.columns, loc="center", cellLoc="center", fontsize=15, colColours=["#05C7F2"]*13)
axes[i//4, i%4].set_title(feature[i])

axes[3, 1].axis('off')
axes[3, 2].axis('off')
axes[3, 3].axis('off')

plt.tight_layout()
plt.show()
```



## Hypothesis Testing

```
In [61]: from scipy.stats import chi2_contingency, mannwhitneyu, chi2
```

H0 : There is no relationship among two variables

H1 : There is relationship among two variables

```
In [62]: hypo_test = {}
nums = ["Age", "Income"]

for i in parquet_df.columns:
    if i in nums:
        mini_df = parquet_df[["Depression", i]]
        cat_one = mini_df[mini_df["Depression"] == "No"][i]
        cat_two = mini_df[mini_df["Depression"] == "Yes"][i]
        test = mannwhitneyu(cat_one, cat_two)
        hypo_test[i] = {"Statistic":test[0], "P-Value":test[1]}
    else:
        test = chi2_contingency(pd.crosstab(parquet_df["Depression"], parquet_df[i]))
        hypo_test[i] = {"Statistic":test[0], "P-Value":test[1]}

hypo_test.pop("Depression")

statistic = [i["Statistic"] for i in hypo_test.values()]
p_Value = [i["P-Value"] for i in hypo_test.values()]
```

```
In [63]: a = pd.DataFrame(data={
    "Target Variable":["Depression (2 Categories)"]*14,
    "Another Variable":["Age (Numerical)", 'Marital Status (4 Categories)', 'Education Level (5 Categories)', 'Number of Children (5 Categories)',
    'Smoking Status (3 Categories)', 'Physical Activity Level (3 Categories)', 'Employment Status (2 Categories)',
    'Income (Numerical)', 'Alcohol Consumption (3 Categories)', 'Dietary Habits (3 Categories)', 'Sleep Patterns (3 Categories)',
    'History of Mental Illness (2 Categories)', 'History of Substance Abuse (2 Categories)',
    'Family History of Depression (2 Categories)'],
    "Type of Hypothesis Test":["Mann-Whitney U Test", 'Chi-Square Test', 'Chi-Square Test', 'Chi-Square Test', 'Chi-Square Test', 'Chi-Square Test',
    'Chi-Square Test', 'Mann-Whitney U Test', 'Chi-Square Test', 'Chi-Square Test', 'Chi-Square Test', 'Chi-Square Test', 'Chi-Square Test', 'Chi-
Square Test'],
    "Statistic":statistic,
    "P-Value":p_Value,
    "Relationship":["There is a relationship" if i<0.05 else "No relationship" for i in p_Value]

})
# a
# print(a.to_markdown())
```

	Target Variable	Another Variable	Type of Hypothesis Test	Statistic	P-Value	Relationship
0	Depression (2 Categories)	Age (Numerical)	Mann-Whitney U Test	1.93121e+10	8.65343e-30	There is a relationship
1	Depression (2 Categories)	Marital Status (4 Categories)	Chi-Square Test	66.9861	1.88552e-14	There is a relationship
2	Depression (2 Categories)	Education Level (5 Categories)	Chi-Square Test	4.92712	0.294857	No relationship
3	Depression (2 Categories)	Number of Children (5 Categories)	Chi-Square Test	37.4444	1.45881e-07	There is a relationship
4	Depression (2 Categories)	Smoking Status (3 Categories)	Chi-Square Test	1166.17	5.8947e-254	There is a relationship
5	Depression (2 Categories)	Physical Activity Level (3 Categories)	Chi-Square Test	527.441	2.93487e-115	There is a relationship
6	Depression (2 Categories)	Employment Status (2 Categories)	Chi-Square Test	744.135	7.56388e-164	There is a relationship
7	Depression (2 Categories)	Income (Numerical)	Mann-Whitney U Test	1.97126e+10	2.13549e-111	There is a relationship
8	Depression (2 Categories)	Alcohol Consumption (3 Categories)	Chi-Square Test	106.272	8.38038e-24	There is a relationship

	Target Variable	Another Variable	Type of Hypothesis Test	Statistic	P-Value	Relationship
9	Depression (2 Categories)	Dietary Habits (3 Categories)	Chi-Square Test	114.2	1.59109e-25	There is a relationship
10	Depression (2 Categories)	Sleep Patterns (3 Categories)	Chi-Square Test	123.653	1.40935e-27	There is a relationship
11	Depression (2 Categories)	History of Mental Illness (2 Categories)	Chi-Square Test	15.0563	0.000104351	There is a relationship
12	Depression (2 Categories)	History of Substance Abuse (2 Categories)	Chi-Square Test	4.50934	0.0337103	There is a relationship
13	Depression (2 Categories)	Family History of Depression (2 Categories)	Chi-Square Test	2.12851	0.144581	No relationship

```
In [64]: df = parquet_df.copy()
df.drop(columns=["Education Level", "Family History of Depression"], inplace=True)
```

```
In [ ]:
```

```
In [65]: df.head()
```

```
Out[65]:
```

	Age	Marital Status	Number of Children	Smoking Status	Physical Activity Level	Employment Status	Income	Alcohol Consumption	Dietary Habits	Sleep Patterns	History of Mental Illness	History of Substance Abuse	Depression
0	31	Married	2	Non-smoker	Active	Unemployed	26265.67	Moderate	Moderate	Fair	Yes	No	Yes
1	55	Married	1	Non-smoker	Sedentary	Employed	42710.36	High	Unhealthy	Fair	Yes	No	Yes
2	78	Widowed	1	Non-smoker	Sedentary	Employed	125332.79	Low	Unhealthy	Good	No	No	No
3	58	Divorced	3	Non-smoker	Moderate	Unemployed	9992.78	Moderate	Moderate	Poor	No	No	No
4	18	Single	0	Non-smoker	Sedentary	Unemployed	8595.08	Low	Moderate	Fair	Yes	No	Yes

## Splitting the data

```
In [66]: le_object = LabelEncoder()
df["Depression"] = le_object.fit_transform(df["Depression"])
raw_df = df
target = df["Depression"]
df = df.drop(columns="Depression")
```

```
In [67]: x_train, x_test, y_train, y_test = train_test_split(df, target, test_size=0.05, random_state=5)
```

```
In [68]: print(f'''
x_train shape : {x_train.shape}
x_test shape : {x_test.shape}
y_train shape : {y_train.shape}
y_test shape : {y_test.shape}
''')
```

```
x_train shape : (393079, 12)
x_test shape : (20689, 12)
y_train shape : (393079,)
y_test shape : (20689,)
```

## Transformation

### Scaling

```
In [69]: x_train.select_dtypes(include="number").head(5)
```

```
Out[69]:
```

	Age	Income
13513	49	104852.57
324041	59	24365.22
144337	78	29327.34
250003	43	27478.55
47584	75	9891.63

```
In [70]: scaler_objects = {}
def do_scaling(x_train, x_test, scaler_object=MinMaxScaler()):

    for i in x_train.select_dtypes(include="number").columns:
        scaler_objects[i] = MinMaxScaler()
        scaler_objects[i].fit(x_train[[i]])
        x_train[i] = scaler_objects[i].transform(x_train[[i]])

    for i in x_test.select_dtypes(include="number").columns:
        x_test[i] = scaler_objects[i].transform(x_test[[i]])

    return x_train, x_test
```

```
In [71]: x_train, x_test = do_scaling(x_train=x_train, x_test=x_test)
```

```
In [72]: x_train.head()
```

```
Out[72]:
```

	Age	Marital Status	Number of Children	Smoking Status	Physical Activity Level	Employment Status	Income	Alcohol Consumption	Dietary Habits	Sleep Patterns	History of Mental Illness	History of Substance Abuse
13513	0.500000	Married	1	Current	Moderate	Employed	0.499308	High	Unhealthy	Poor	No	No
324041	0.661290	Married	1	Non-smoker	Sedentary	Employed	0.116026	Moderate	Moderate	Fair	Yes	Yes
144337	0.967742	Married	3	Non-smoker	Moderate	Employed	0.139655	Moderate	Moderate	Fair	No	Yes
250003	0.403226	Married	1	Former	Sedentary	Unemployed	0.130852	High	Moderate	Poor	Yes	No
47584	0.919355	Widowed	2	Former	Moderate	Unemployed	0.047102	High	Unhealthy	Poor	Yes	No

```
In [73]: x_test.head()
```

	Age	Marital Status	Number of Children	Smoking Status	Physical Activity Level	Employment Status	Income	Alcohol Consumption	Dietary Habits	Sleep Patterns	History of Mental Illness	History of Substance Abuse
366559	0.080645	Married	2	Non-smoker	Active	Unemployed	0.138907	Moderate	Healthy	Fair	No	No
66432	0.129032	Single	0	Current	Sedentary	Employed	0.354114	High	Unhealthy	Fair	Yes	Yes
243274	0.935484	Married	2	Non-smoker	Moderate	Employed	0.613201	High	Healthy	Poor	No	No
209153	0.096774	Married	1	Non-smoker	Active	Employed	0.410786	Low	Healthy	Poor	No	No
394294	0.919355	Married	1	Non-smoker	Moderate	Employed	0.416343	Moderate	Healthy	Fair	No	No

```
In [74]: # scaler_objects = {}

# for i in x_train.select_dtypes(include="number").columns:
#     scaler_objects[i] = MinMaxScaler()
#     scaler_objects[i].fit(x_train[[i]])
#     x_train[i] = scaler_objects[i].transform(x_train[[i]])

In [75]: x_train.select_dtypes(include="number").head(5)
```

	Age	Income
13513	0.500000	0.499308
324041	0.661290	0.116026
144337	0.967742	0.139655
250003	0.403226	0.130852
47584	0.919355	0.047102

```
In [76]: for i in x_test.select_dtypes(include="number").columns:
x_test[i] = scaler_objects[i].transform(x_test[[i]])
```

## Encoding

```
In [77]: new_df = x_train.select_dtypes(include="number")
new_df.reset_index(inplace=True)
encoder_objects = {}

for i in x_train.select_dtypes(exclude="number").columns:
    # print(i)
    encoder_objects[i] = OneHotEncoder(dtype='int', drop="first")
    dummy_df = encoder_objects[i].fit_transform(x_train[[i]]).toarray()
    new_df = pd.concat([new_df, pd.DataFrame(dummy_df, columns=encoder_objects[i].get_feature_names_out())], axis=1)
    # print(new_df)

x_train = new_df

In [78]: new_df = x_test.select_dtypes(include="number")
new_df.reset_index(inplace=True)

for i in x_test.select_dtypes(exclude="number").columns:
    dummy_df = encoder_objects[i].transform(x_test[[i]]).toarray()
    new_df = pd.concat([new_df, pd.DataFrame(dummy_df, columns=encoder_objects[i].get_feature_names_out())], axis=1)
    # print(new_df)

x_test = new_df

In [79]: x_train.drop(columns="index", inplace=True)
x_test.drop(columns="index", inplace=True)
```

## Base Model - Logistic Regression

```
In [80]: def get_metrics(model, x_test, y_test):
...
    Parameters:
    -----
    model : ML Model Object i.e., LinearRegression()

    x_test : Features Testing DataFrame from train_test_split()

    y_test : Target Testing DataFrame from train_test_split()

    Returns:
    -----
    subplots(3) : Clasification Report, Confusion Matrix, ROC-AUC Curve

    ...

    y_pred = model.predict(x_test)
    fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(20, 8))

    model_name = model if len(str(model)) <= 50 else str(model)[:50]+"..."
    fig.suptitle(f"\nClassification Metrics for {model_name}\n", fontsize=30)

    #####
    cr = pd.DataFrame(classification_report(y_test, y_pred, output_dict=True)).T.iloc[:, :3]
    accuracy = round((float(cr.loc["accuracy"][0]))*100, 2)
    t = cr.iloc[[0,1,3,4], :]
    t.loc[["macro avg", "weighted avg"]] = t.loc[["macro avg", "weighted avg"]]*100
    t = round(t, 2)

    axes[0].axis('tight')
    axes[0].axis('off')
    t = axes[0].table(cellText=t.values, colLabels=t.columns, rowLabels=t.index, loc="center", cellloc="center", rowLoc='right', fontsize=5, colColours=["#05C7F2"]*3,
rowColours=["#BCB7B1"]*4, colWidths=[0.25]*3)
    t.scale(1, 3)
    axes[0].set_title(f"Classification Report\n", fontsize=15)
    axes[0].text(x=0.018, y=0.035, s=f"Accuracy : {accuracy}%", fontsize=15)
```

```

#####
#####
#####
l = np.array(["True Negative", "False Positive", "False Negative", "True Positive"])
cm = confusion_matrix(y_test, y_pred)
c = cm.flatten()
sns.heatmap(confusion_matrix(y_test, y_pred), annot=np.array([f"{i}\n{j}," for i, j in zip(l, c)]).reshape(2,2), fmt="", cbar=False, ax=axes[1], cmap =
ListedColormap(['#05C7F2']), linewidths = 1, annot_kws = {'size':15})
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        if i != j: # Check if not diagonal
            # Set cell background to white
            rect = plt.Rectangle((j, i), 1, 1, facecolor='#8CB7B1')
            axes[1].add_patch(rect)
axes[1].set_xlabel("Predicted Values")
axes[1].set_ylabel("Actual Values")
axes[1].axis('tight')
axes[1].set_title(f"Confusion Matrix\n", fontsize=15)
#####
#####
fpr, tpr, _ = roc_curve(y_test, model.predict_proba(x_test)[:,-1])
axes[2].plot(fpr, tpr, color='#05C7F2', label=f"ROC-AUC Score\n{roc_auc_score(y_test, model.predict_proba(x_test)[:,-1]):.4f}", fillstyle="full")
axes[2].plot([0, 1], [0, 1], color='#A19C95', linestyle='dashed')
axes[2].fill_between(fpr, tpr, alpha=0.5, color='#8CB7B1')
axes[2].set_xlabel("False Positive Rate")
axes[2].set_ylabel("True Positive Rate")
axes[2].legend(loc="upper left")
axes[2].set_title(f"ROC-AUC Curve\n", fontsize=15)
#####

plt.tight_layout()
plt.show()


```

```

In [81]: model = LogisticRegression().fit(x_train, y_train)
get_metrics(model, x_test, y_test)

```

## Classification Metrics for LogisticRegression()



```

In [82]: from sklearn.tree import DecisionTreeClassifier


```

```

In [83]: model = DecisionTreeClassifier().fit(x_train, y_train)
get_metrics(model, x_test, y_test)

```

## Classification Metrics for DecisionTreeClassifier()



```

In [84]: from sklearn.ensemble import RandomForestClassifier


```

```

In [85]: model = RandomForestClassifier().fit(x_train, y_train)
get_metrics(model, x_test, y_test)

```


## Classification Metrics for RandomForestClassifier()




```
In [86]: from xgboost.sklearn import XGBClassifier
```

```
In [87]: model = XGBClassifier().fit(x_train, y_train)
get_metrics(model, x_test, y_test)
```

## Classification Metrics for XGBClassifier(base\_score=None, booster=None, callb...)




```
In [88]: plt.figure(figsize=(20,5))
sns.scatterplot(x=row_df['Age'], y=row_df['Income'], hue=row_df['Depression'], palette=["Black", "Orange"])
plt.tight_layout()
plt.show()
```



```
In [89]: plt.figure(figsize=(20,5))
sns.lineplot(row_df[row_df["Depression"] == 0]["Age"].value_counts(), color="Black", errorbar=None, label="Depression - 0")
sns.lineplot(row_df[row_df["Depression"] == 1]["Age"].value_counts(), color="Orange", errorbar=None, label="Depression - 1")
plt.legend(loc="center right")
plt.tight_layout()
plt.show()
```





```
In [90]: dummy_df = parquet_df.select_dtypes(include="number")
```

```
In [91]: for i in parquet_df.select_dtypes(exclude="number").columns:
dummy_df[i] = LabelEncoder().fit_transform(parquet_df[i])
```

```
In [92]: # a = OneHotEncoder(drop="first").fit(parquet_df.select_dtypes(exclude="number"))
# b = a.transform(parquet_df.select_dtypes(exclude="number")).toarray()
# for i in range(len(a.get_feature_names_out())):
#     dummy_df[a.get_feature_names_out()[i]] = b[:, i]
```

```
In [ ]:
```

```
In [93]: dummy_df
```

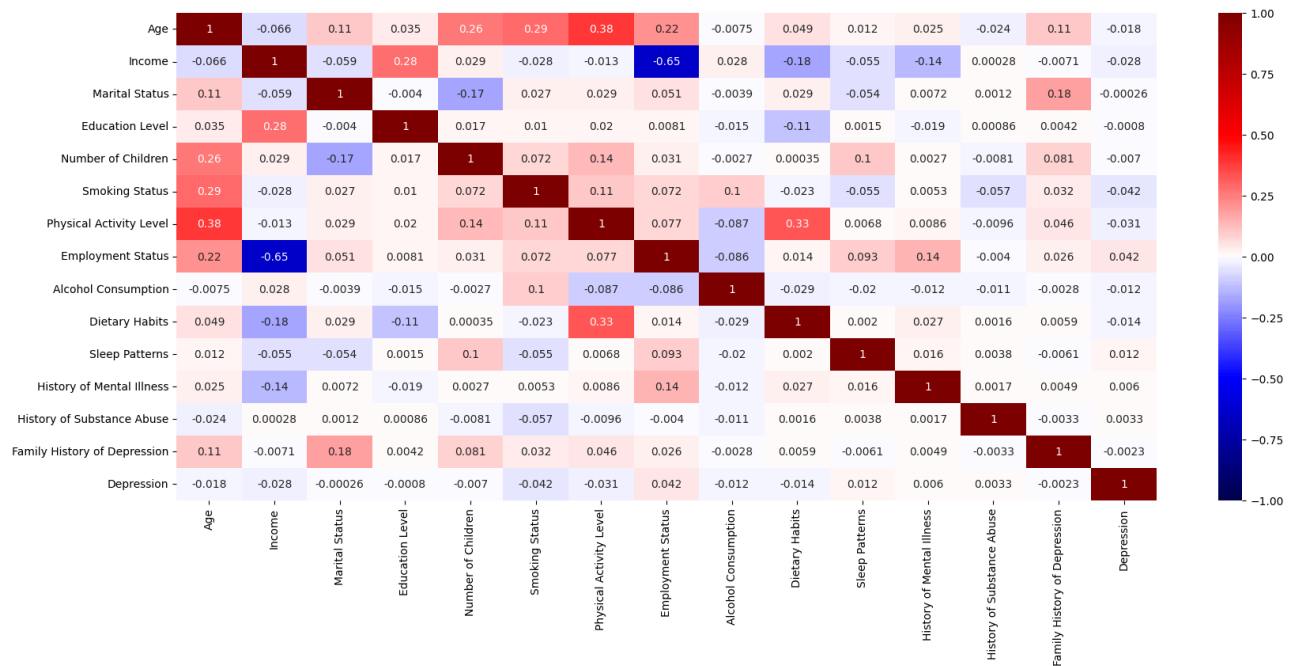
```
Out[93]:
```

	Age	Income	Marital Status	Education Level	Number of Children	Smoking Status	Physical Activity Level	Employment Status	Alcohol Consumption	Dietary Habits	Sleep Patterns	History of Mental Illness	History of Substance Abuse	Family History of Depression	Depression
0	31	26265.67	1	1	2	2	0	1	2	1	0	1	0	1	1
1	55	42710.36	1	2	1	2	2	0	0	2	0	1	0	0	1
2	78	125332.79	3	3	1	2	2	0	1	2	1	0	0	1	0
3	58	9992.78	0	3	3	2	1	1	2	1	2	0	0	0	0
4	18	8595.08	2	2	0	2	2	1	1	1	0	1	0	1	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
413763	68	109233.43	1	3	0	1	1	0	1	0	1	0	0	0	0
413764	26	96760.97	2	1	0	0	0	0	1	0	2	1	1	0	1
413765	57	77353.26	1	1	0	2	2	0	2	1	0	0	0	1	1
413766	71	24557.08	1	0	2	2	2	1	2	1	2	0	1	0	0
413767	62	107125.74	3	3	0	1	1	0	2	0	1	0	1	0	0

413768 rows x 15 columns

```
In [94]: plt.figure(figsize=(20,8))
sns.heatmap(dummy_df.corr(), vmin=-1, annot=True, cmap="seismic")
```

```
Out[94]: <Axes: >
```



## Logistic Regression with class weights

```
In [95]: model = LogisticRegression(class_weight={0:1, 1:1.3}).fit(x_train, y_train)
get_metrics(model, x_test, y_test)
```


## Classification Metrics for LogisticRegression(class\_weight={0: 1, 1: 1.3})

Classification Report


Accuracy : 66.86%

	precision	recall	f1-score
0	0.67	0.98	0.8
1	0.43	0.03	0.05
macro avg	55.16	50.51	42.67
weighted avg	59.37	66.86	55.46

Confusion Matrix



ROC-AUC Curve



```
In [96]: model = LogisticRegression(class_weight={0:1, 1:2}).fit(x_train, y_train)
get_metrics(model, x_test, y_test)
```


## Classification Metrics for LogisticRegression(class\_weight={0: 1, 1: 2})

Classification Report


Accuracy : 58.26%

	precision	recall	f1-score
0	0.7	0.65	0.68
1	0.38	0.44	0.41
macro avg	54.23	54.5	54.23
weighted avg	59.77	58.26	58.89

Confusion Matrix



ROC-AUC Curve



## Voting Classifier

```
In [97]: train_df = x_train.copy()
train_df["Depression"] = pd.DataFrame(y_train).reset_index(drop=True)["Depression"]
```

```
In [98]: D_1_df = train_df[train_df["Depression"] == 1]
```

```
In [99]: D_1_df.head(3)
```

```
Out[99]:
```

	Age	Income	Marital Status_Married	Marital Status_Single	Marital Status_Widowed	Number of Children_1	Number of Children_2	Number of Children_3	Number of Children_4	Smoking Status_Former	...	Employment Status_Unemployed	Alcohol Consumption_Low	Alcohol Consumption_High
4	0.919355	0.047102	0	0	1	0	1	0	0	1	...	1	0	0
6	0.225806	0.406555	1	0	0	0	0	1	0	0	...	0	0	0
7	0.854839	0.067561	0	0	1	0	0	0	0	1	...	1	0	0

3 rows × 23 columns

```
In [100]: D_0_df = train_df[train_df["Depression"] == 0]
```

```
In [101]: D_0_df.head(3)
```

```
Out[101]:
```

	Age	Income	Marital Status_Married	Marital Status_Single	Marital Status_Widowed	Number of Children_1	Number of Children_2	Number of Children_3	Number of Children_4	Smoking Status_Former	...	Employment Status_Unemployed	Alcohol Consumption_Low	Alcohol Consumption_High
0	0.500000	0.499308	1	0	0	1	0	0	0	0	...	0	0	0
1	0.661290	0.116026	1	0	0	1	0	0	0	0	...	0	0	0
2	0.967742	0.139655	1	0	0	0	0	1	0	0	...	0	0	0

3 rows × 23 columns

```
In [102]: D_1_df.shape, D_0_df.shape
```

```
Out[102]: ((129416, 23), (263663, 23))
```

```
In [103]: predicts = pd.DataFrame({"Actual":y_test})
models = pd.DataFrame()
list_of_models = dict()

for i in range(11):
```

```

D_1 = D_1_df.sample(frac=0.6, random_state=np.random.randint(1, 45781236))
D_0 = D_0_df.sample(n=D_1.shape[0], random_state=np.random.randint(1, 45781236))
mini_data = pd.concat([D_0, D_1]).sample(frac=1)
mini_y_train = mini_data["Depression"]
mini_x_train = mini_data.drop(columns="Depression")
mini_model = XGBClassifier().fit(mini_x_train, mini_y_train)
list_of_models[f"Model {i+1}"] = mini_model
predicts[f"Model {i+1} Prediction"] = mini_model.predict(x_test)
mets = precision_recall_fscore_support(y_test, mini_model.predict(x_test))
models = models._append(
    {
        "Model": f"D_0_df{i+1}",
        "Precision_0": mets[0][0],
        "Precision_1": mets[0][1],
        "Recall_0": mets[1][0],
        "Recall_1": mets[1][1],
        "Accuracy": accuracy_score(y_test, mini_model.predict(x_test))
    }, ignore_index=True)
print(f".....Done Model {i+1}")

```

```

.....Done Model 1
.....Done Model 2
.....Done Model 3
.....Done Model 4
.....Done Model 5
.....Done Model 6
.....Done Model 7
.....Done Model 8
.....Done Model 9
.....Done Model 10
.....Done Model 11

```

In [104..

predicts.head(3)

Out[104..

	Actual	Model 1 Prediction	Model 2 Prediction	Model 3 Prediction	Model 4 Prediction	Model 5 Prediction	Model 6 Prediction	Model 7 Prediction	Model 8 Prediction	Model 9 Prediction	Model 10 Prediction	Model 11 Prediction
366559	0	1	0	1	1	1	0	0	1	1	1	1
66432	1	0	1	1	1	1	1	1	1	1	1	1
243274	0	0	0	1	0	1	0	0	1	1	1	0

In [105..

models

Out[105..

	Model	Precision_0	Precision_1	Recall_0	Recall_1	Accuracy
0	D_0_df1	0.677169	0.333492	0.496402	0.515683	0.502731
1	D_0_df2	0.677177	0.334957	0.557850	0.455750	0.524337
2	D_0_df3	0.678361	0.330788	0.281048	0.727286	0.427522
3	D_0_df4	0.679058	0.332516	0.373291	0.638934	0.460486
4	D_0_df5	0.666220	0.323353	0.465031	0.523192	0.484122
5	D_0_df6	0.677965	0.335712	0.551230	0.464144	0.522645
6	D_0_df7	0.667907	0.326042	0.361491	0.632160	0.450336
7	D_0_df8	0.687079	0.338231	0.403655	0.623767	0.475905
8	D_0_df9	0.680710	0.330133	0.176716	0.830364	0.391271
9	D_0_df10	0.668586	0.327267	0.233991	0.762627	0.407511
10	D_0_df11	0.677725	0.329236	0.144050	0.859814	0.378994

In [106..

```

model = VotingClassifier(estimators=list(list_of_models.items()), voting="soft", verbose=True).fit(x_train, y_train)

```

```

[Voting] ..... (1 of 11) Processing Model 1, total= 0.6s
[Voting] ..... (2 of 11) Processing Model 2, total= 0.6s
[Voting] ..... (3 of 11) Processing Model 3, total= 0.6s
[Voting] ..... (4 of 11) Processing Model 4, total= 0.6s
[Voting] ..... (5 of 11) Processing Model 5, total= 0.6s
[Voting] ..... (6 of 11) Processing Model 6, total= 0.6s
[Voting] ..... (7 of 11) Processing Model 7, total= 0.7s
[Voting] ..... (8 of 11) Processing Model 8, total= 0.6s
[Voting] ..... (9 of 11) Processing Model 9, total= 0.6s
[Voting] ..... (10 of 11) Processing Model 10, total= 0.6s
[Voting] ..... (11 of 11) Processing Model 11, total= 0.7s

```

In [107..

```

get_metrics(model, x_test, y_test)

```

## Classification Metrics for VotingClassifier(estimators=[('Model 1', ...)

Classification Report

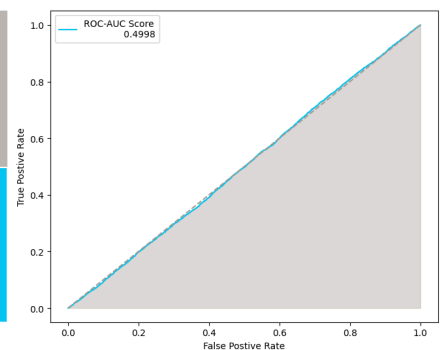
Accuracy : 56.51%

	precision	recall	f1-score
0	0.67	0.69	0.68
1	0.33	0.3	0.31
macro avg	49.78	49.79	49.75
weighted avg	55.71	56.51	56.08

Confusion Matrix

	Actual Values	0	1
Predicted Values	0	True Negative 9,639	False Positive 4,259
	1	False Negative 4,739	True Positive 2,052

ROC-AUC Curve



```

In [108... parquet_df["Depression"].value_counts()

Out[108...
Depression
No    277561
Yes   136207
Name: count, dtype: int64

In [109... from imblearn.over_sampling import SMOTE


In [110... x_train_SMOTE, y_train_SMOTE = SMOTE(sampling_strategy={1:240451}, random_state=30, k_neighbors=1).fit_resample(x_train, y_train)

In [111... model_SMOTE = XGBClassifier().fit(x_train_SMOTE, y_train_SMOTE)

In [112... get_metrics(model_SMOTE, x_test, y_test)

```

## Classification Metrics for XGBClassifier(base\_score=None, booster=None, callb...)



## Seperate Models for Numeric & Categorical

```

In [113... num_df = parquet_df.select_dtypes(include='number')

In [114... num_df.head()

Out[114...
   Age  Income
0   31  26265.67
1   55  42710.36
2   78 125332.79
3   58   9992.78
4   18   8595.08

In [ ]: x_train, x_test, y_train, y_test = train_test_split(num_df.drop(columns="Depression"), num_df["Depression"], test_size=0.05, random_state=5)

In [ ]: print(f'''
x_train shape : {x_train.shape}
x_test shape : {x_test.shape}
y_train shape : {y_train.shape}
y_test shape : {y_test.shape}
...''')

In [ ]: x_test.head()

In [ ]: x_train, x_test = do_scaling(x_train, x_test)

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]: cat_df = parquet_df.select_dtypes(exclude='number')
cat_df = pd.concat([cat_df, parquet_df["Depression"]], axis=1)

In [ ]: cat_df.head()

In [ ]:

```